


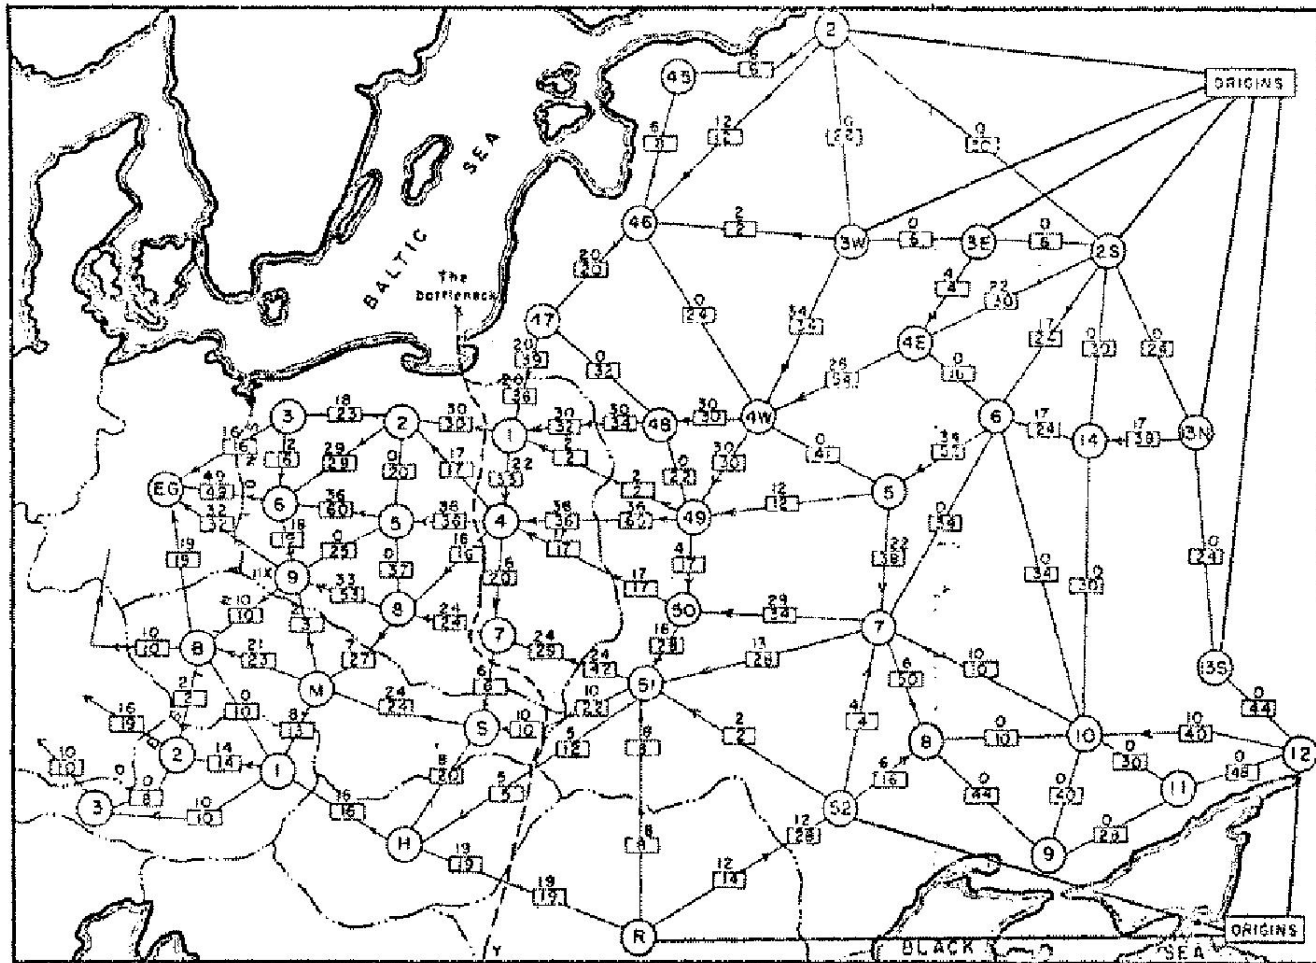
Chapter 7

Network Flow



Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.

Soviet Rail Network, 1955



Reference: *On the history of the transportation and maximum flow problems.*
Alexander Schrijver in Math Programming, 91: 3, 2002.

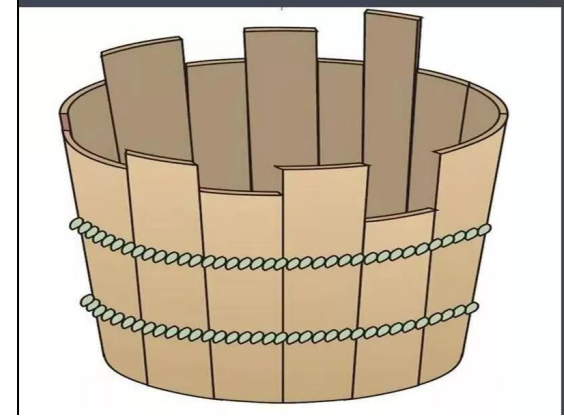
Maximum Flow and Minimum Cut

Max flow and min cut.

- Two very rich algorithmic problems.
- Cornerstone problems in combinatorial optimization.
- Beautiful mathematical duality.

Nontrivial applications / reductions.

- Data mining.
- Open-pit mining.
- Project selection.
- Airline scheduling.
- Bipartite matching.
- Baseball elimination.
- Image segmentation.
- Network connectivity.
- Network reliability.
- Distributed computing.
- Egalitarian stable matching.
- Security of statistical data.
- Network intrusion detection.
- Multi-camera scene reconstruction.
- Many many more ...



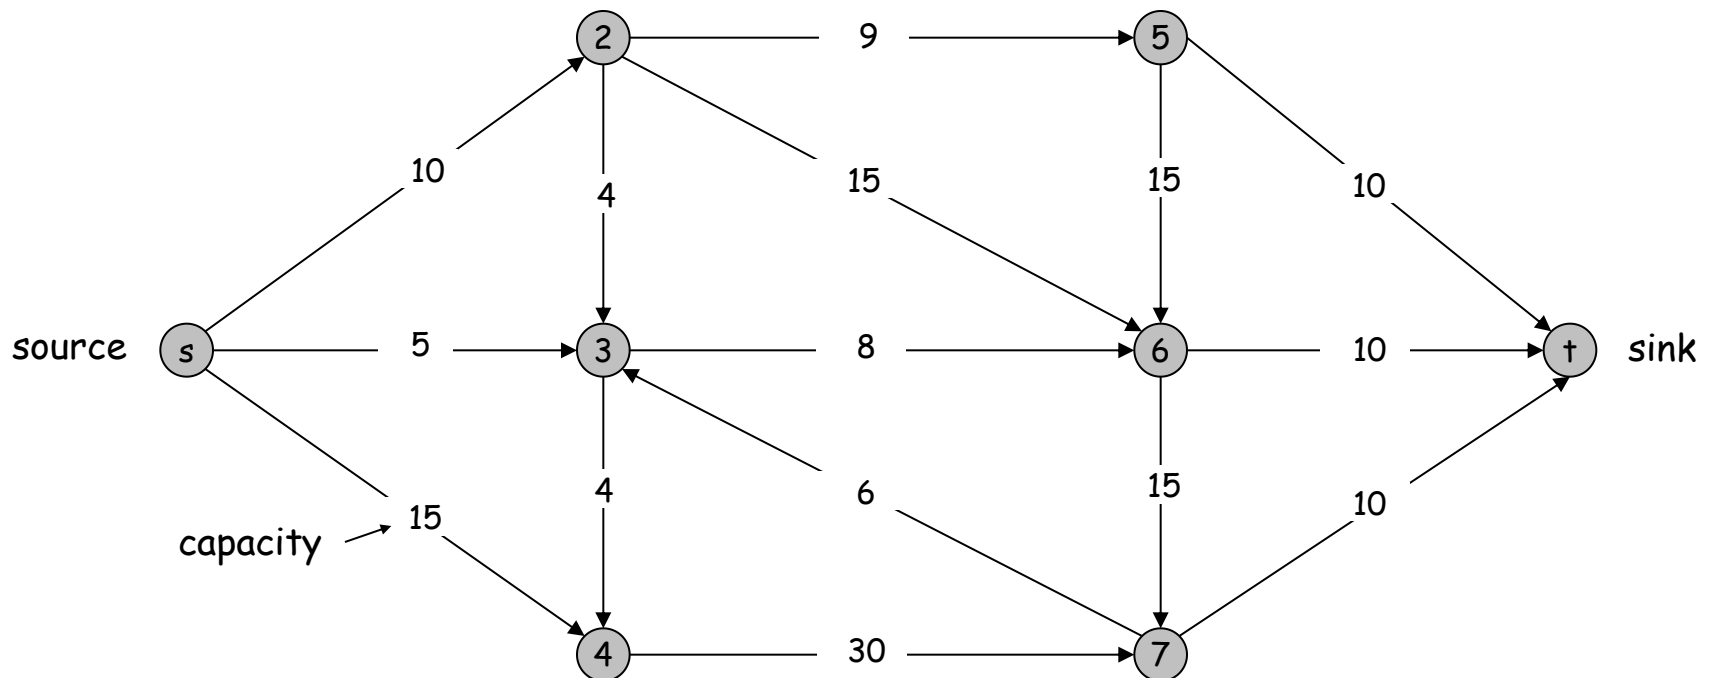
the barrel effect

the capacity of a barrel is determined not by the longest wooden bars, but by the shortest

Minimum Cut Problem

Flow network.

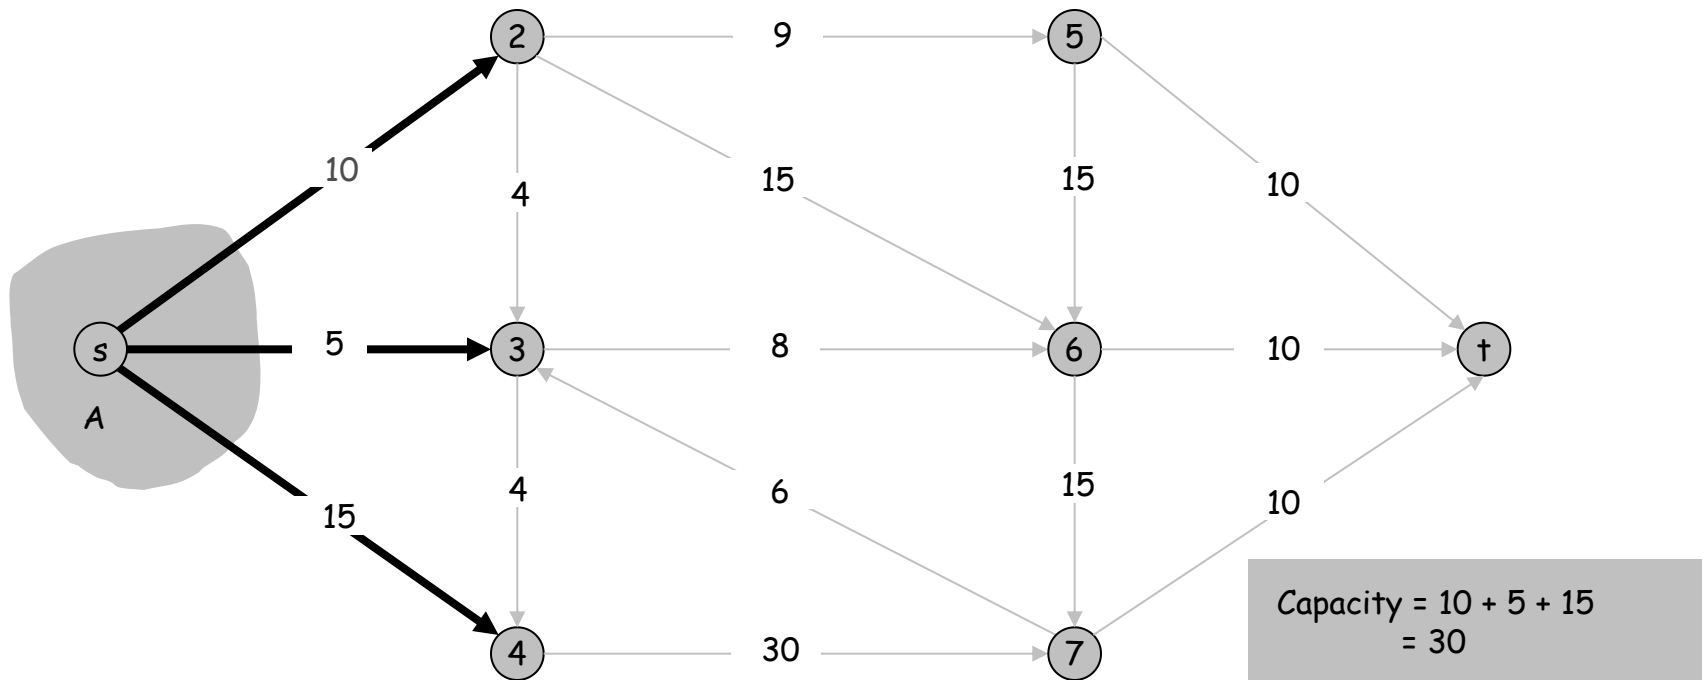
- Abstraction for material **flowing** through the edges.
- $G = (V, E)$ = directed graph, no parallel edges.
- Two distinguished nodes: s = source, t = sink.
- $c(e)$ = capacity of edge e .



Cuts

Def. An **s-t cut** is a partition (A, B) of V with $s \in A$ and $t \in B$.


Def. The **capacity** of a cut (A, B) is: $cap(A, B) = \sum_{e \text{ out of } A} c(e)$



Cuts


Def. An **s-t cut** is a partition (A, B) of V with $s \in A$ and $t \in B$.

Def. The **capacity** of a cut (A, B) is: $cap(A, B) = \sum_{e \text{ out of } A} c(e)$



Minimum Cut Problem

Min s-t cut problem. Find an s-t cut of minimum capacity.

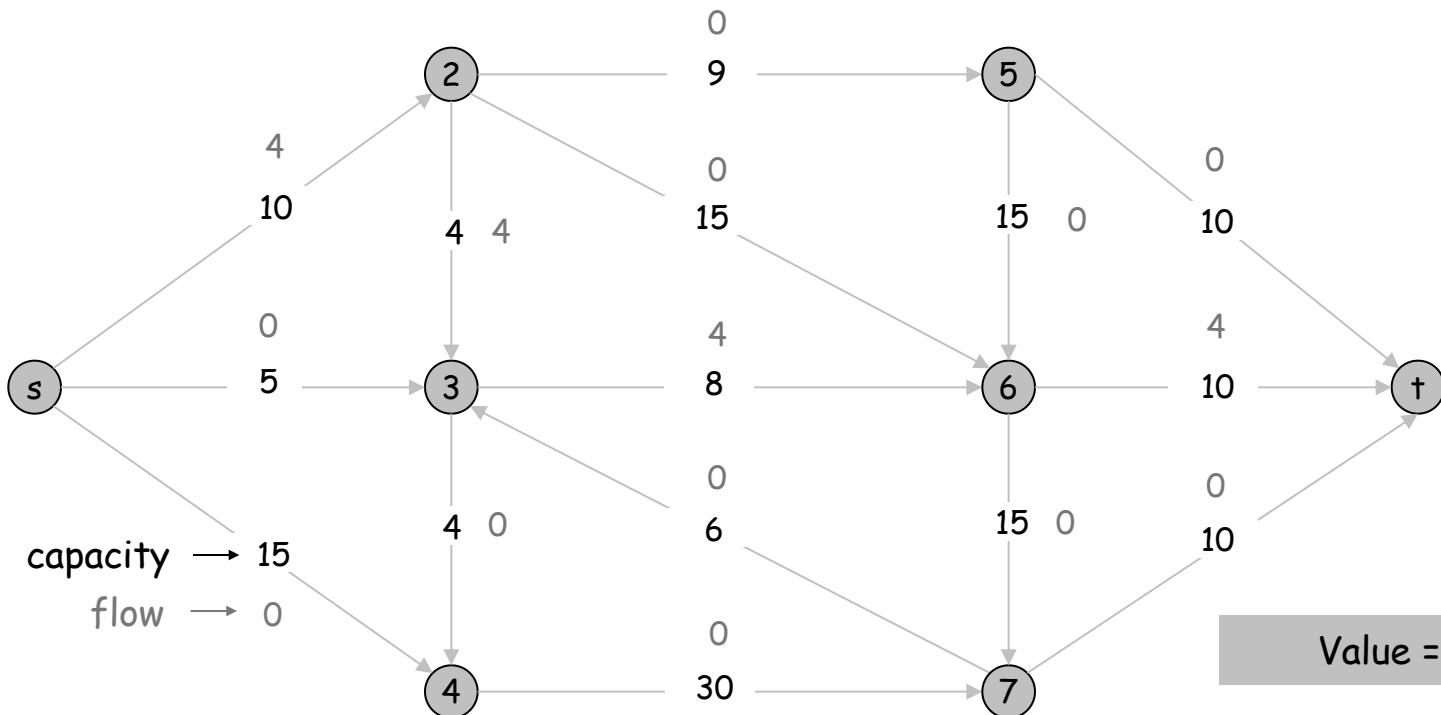


Flows

Def. An **s-t flow** is a function that satisfies:

- For each $e \in E$: $0 \leq f(e) \leq c(e)$ [capacity]
- For each $v \in V - \{s, t\}$: $\sum_{e \text{ in to } v} f(e) = \sum_{e \text{ out of } v} f(e)$ [conservation]

Def. The **value** of a flow f is: $v(f) = \sum_{e \text{ out of } s} f(e)$.

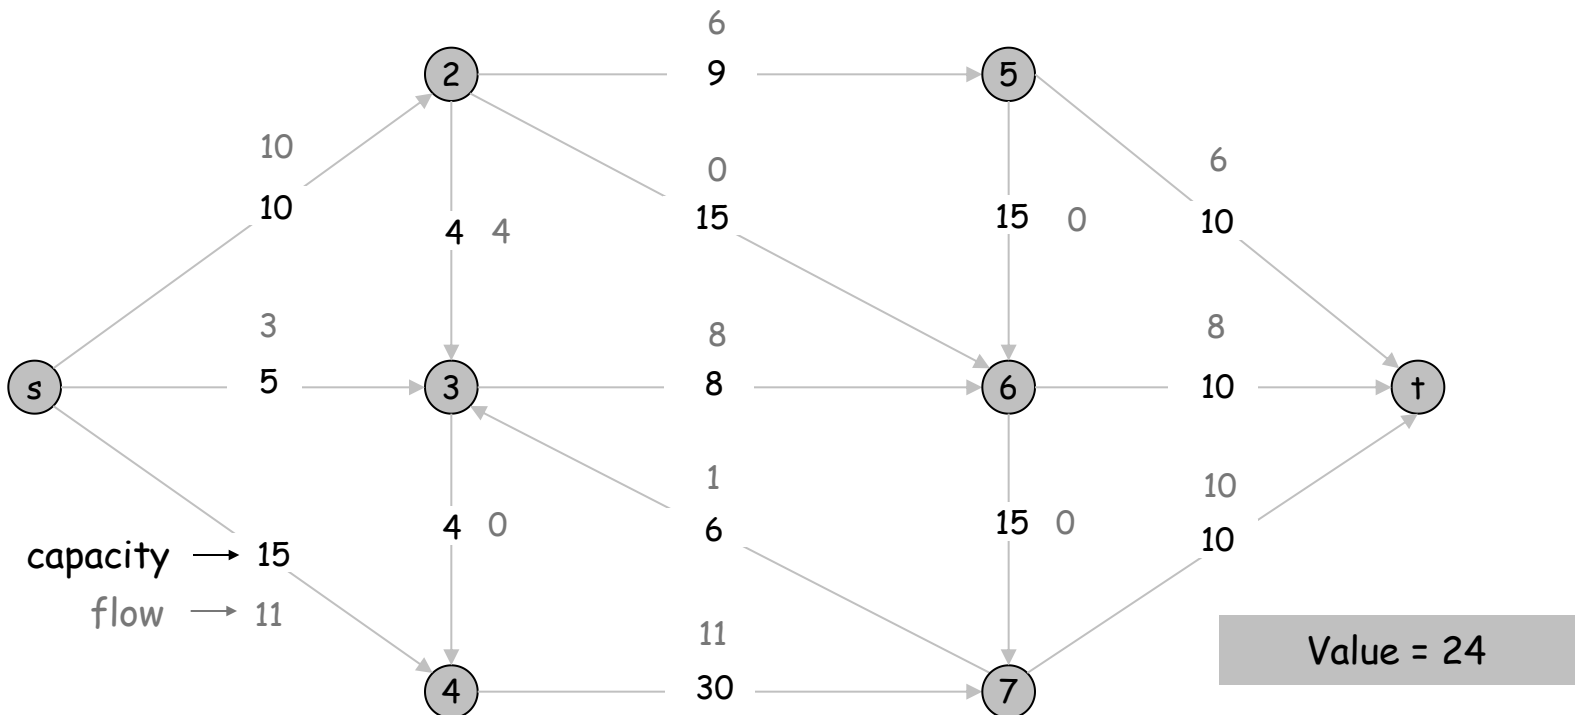


Flows

Def. An **s-t flow** is a function that satisfies:

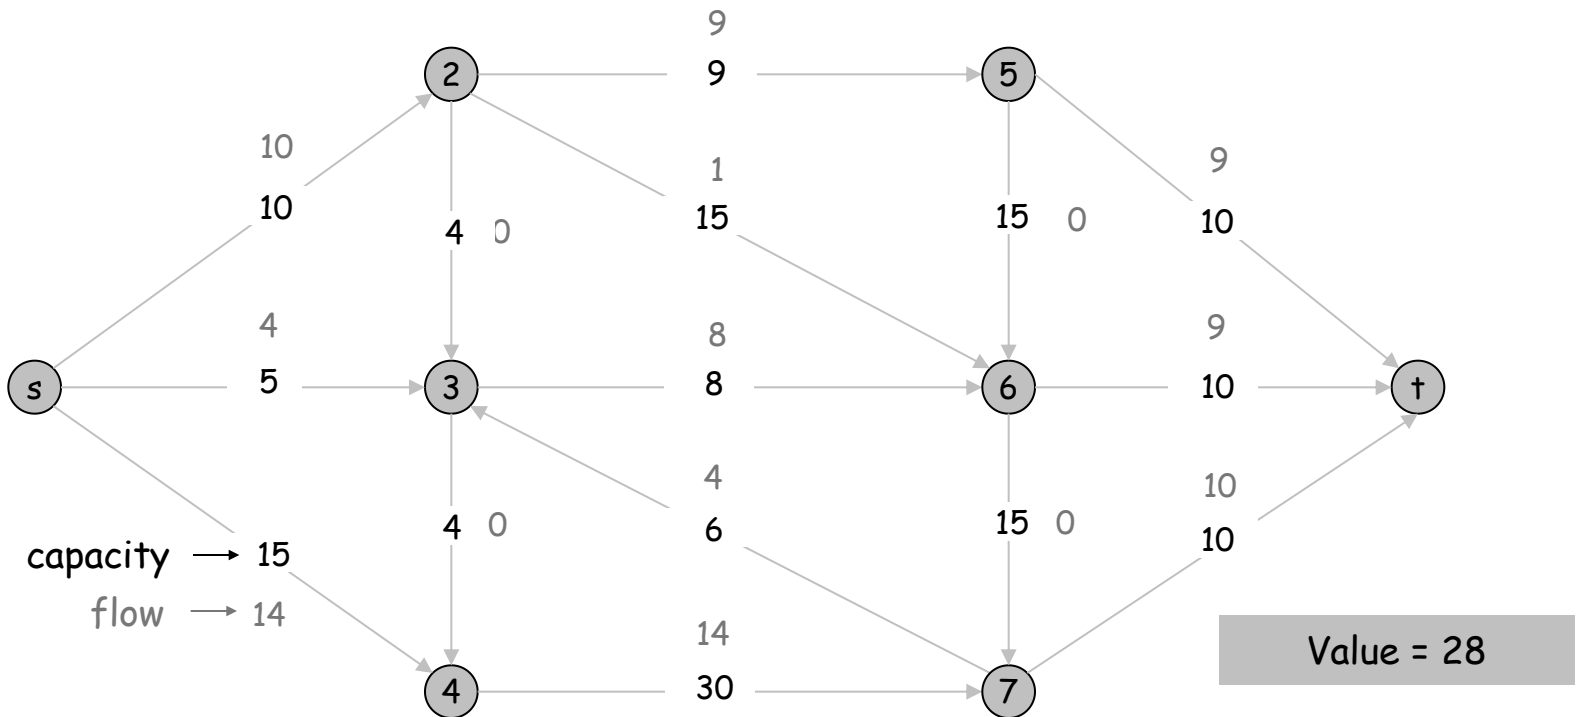
- For each $e \in E$: $0 \leq f(e) \leq c(e)$ [capacity]
- For each $v \in V - \{s, t\}$: $\sum_{e \text{ in to } v} f(e) = \sum_{e \text{ out of } v} f(e)$ [conservation]

Def. The **value** of a flow f is: $v(f) = \sum_{e \text{ out of } s} f(e)$.



Maximum Flow Problem

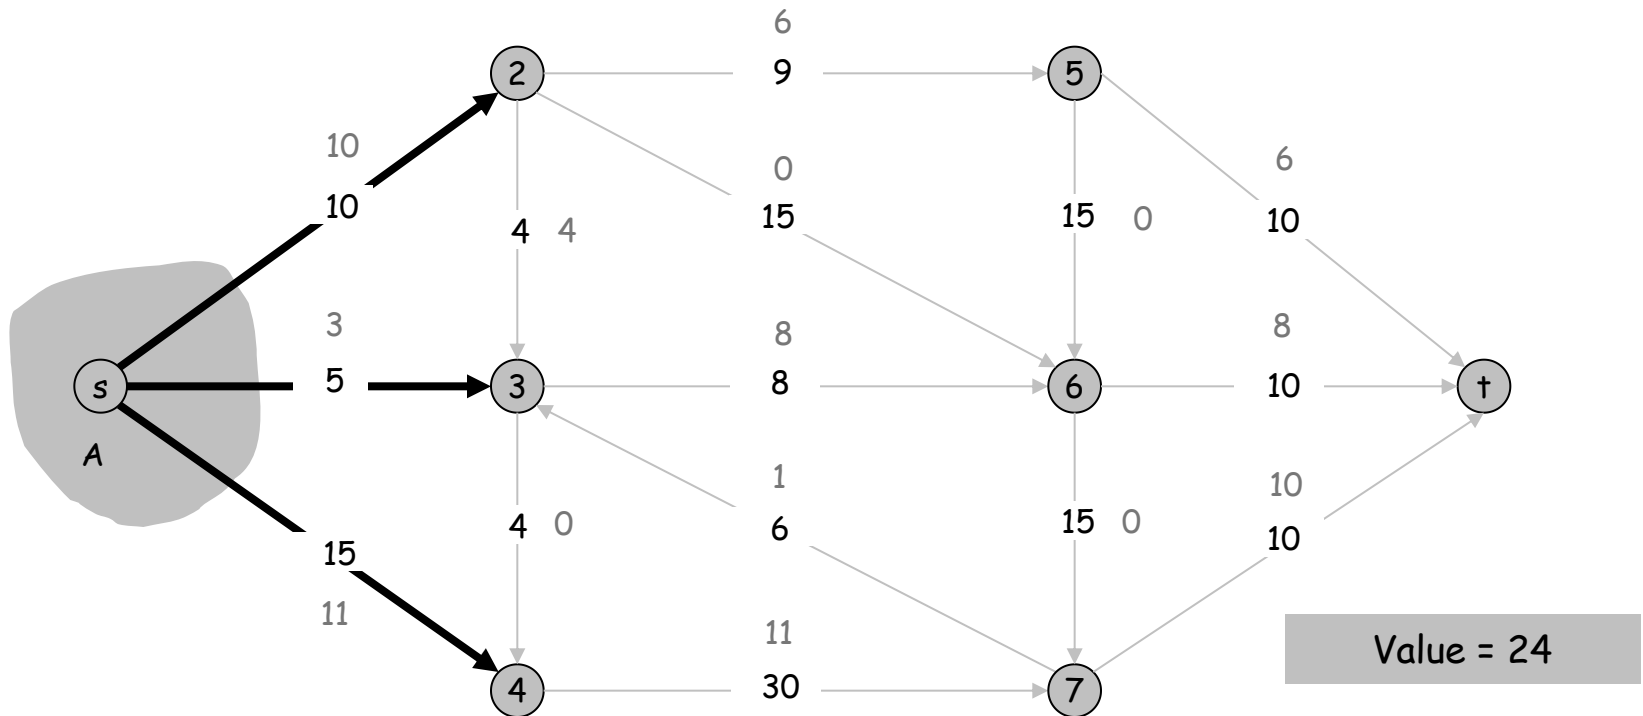
Max flow problem. Find s-t flow of maximum value.



Flows and Cuts

Flow value lemma. Let f be any flow, and let (A, B) be any s - t cut. Then, the net flow sent across the cut is equal to the amount leaving s .


$$\sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) = v(f)$$



Flows and Cuts

Flow value lemma. Let f be any flow, and let (A, B) be any s - t cut. Then, the net flow sent across the cut is equal to the amount leaving s .

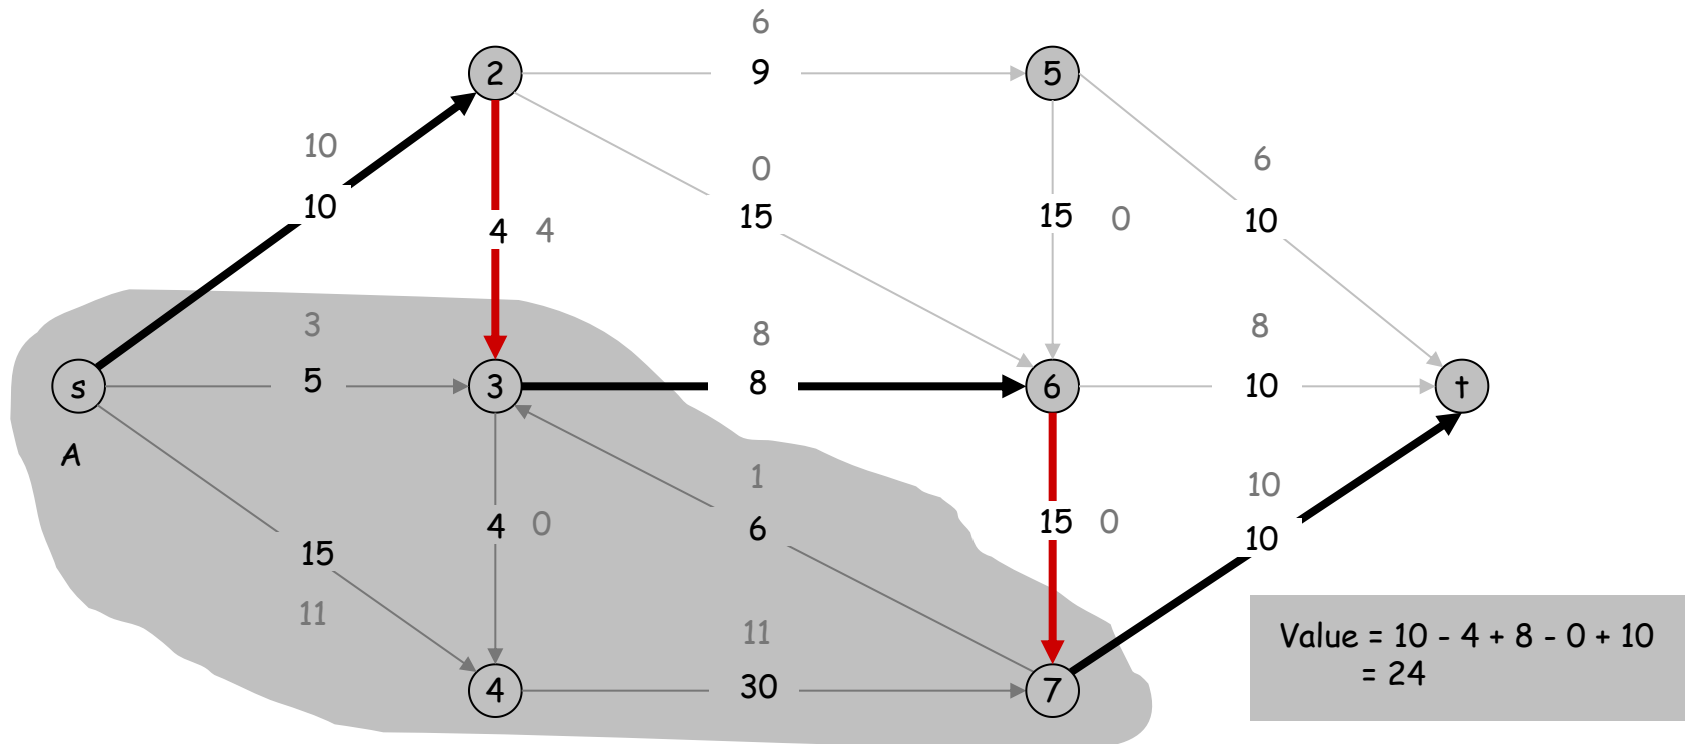
$$\sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) = v(f)$$



Flows and Cuts

Flow value lemma. Let f be any flow, and let (A, B) be any s - t cut. Then, the net flow sent across the cut is equal to the amount leaving s .

$$\sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) = v(f)$$



Flows and Cuts

Flow value lemma. Let f be any flow, and let (A, B) be any s - t cut. Then, the net flow sent across the cut is equal to the amount leaving s .

$$\sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) = v(f).$$

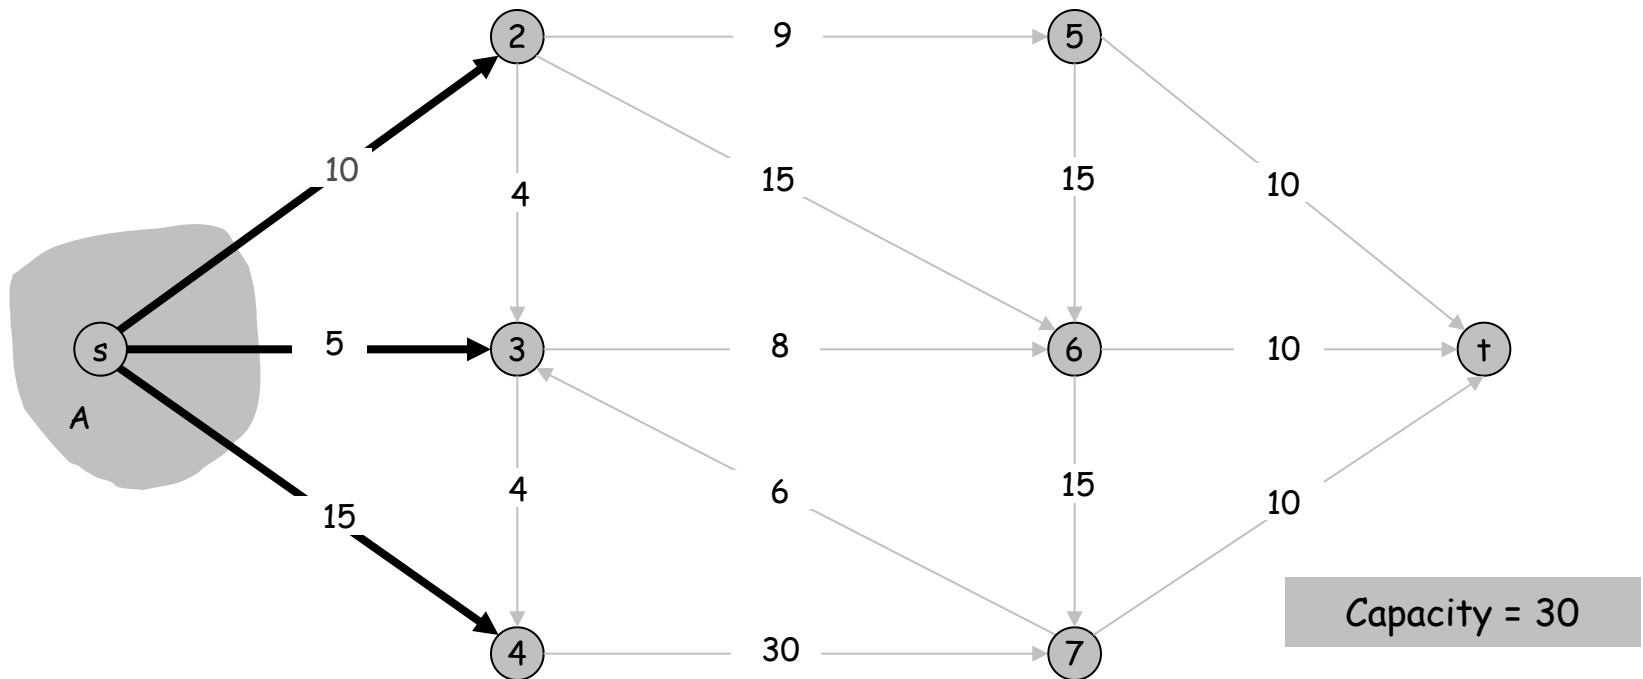
Pf.

$$\begin{aligned} v(f) &= \sum_{e \text{ out of } s} f(e) \\ \text{by flow conservation, all terms} &\longrightarrow = \sum_{v \in A} \left(\sum_{e \text{ out of } v} f(e) - \sum_{e \text{ in to } v} f(e) \right) \\ \text{except } v = s \text{ are } 0 & \\ &= \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e). \end{aligned}$$

Flows and Cuts

Weak duality. Let f be any flow, and let (A, B) be any s - t cut. Then the value of the flow is at most the capacity of the cut.

Cut capacity = 30 \Rightarrow Flow value ≤ 30

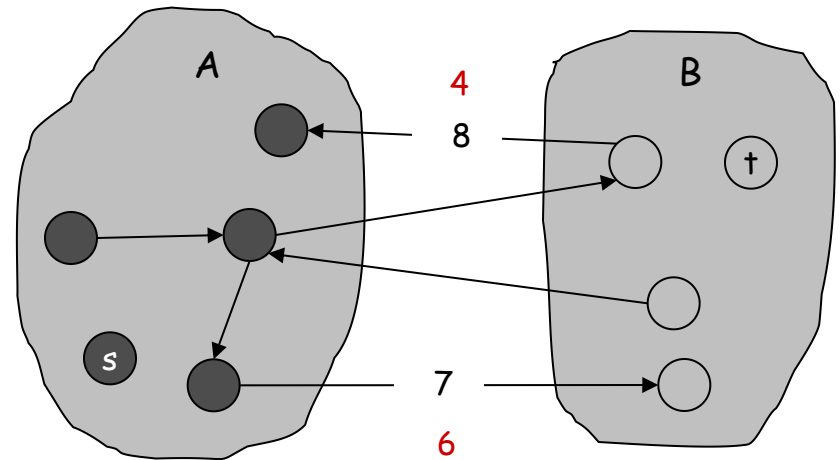


Flows and Cuts

Weak duality. Let f be any flow. Then, for any s - t cut (A, B) we have $v(f) \leq \text{cap}(A, B)$.

Pf.

$$\begin{aligned} v(f) &= \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) \\ &\leq \sum_{e \text{ out of } A} c(e) \\ &\leq \sum_{e \text{ out of } A} c(e) \\ &= \text{cap}(A, B) \quad \blacksquare \end{aligned}$$

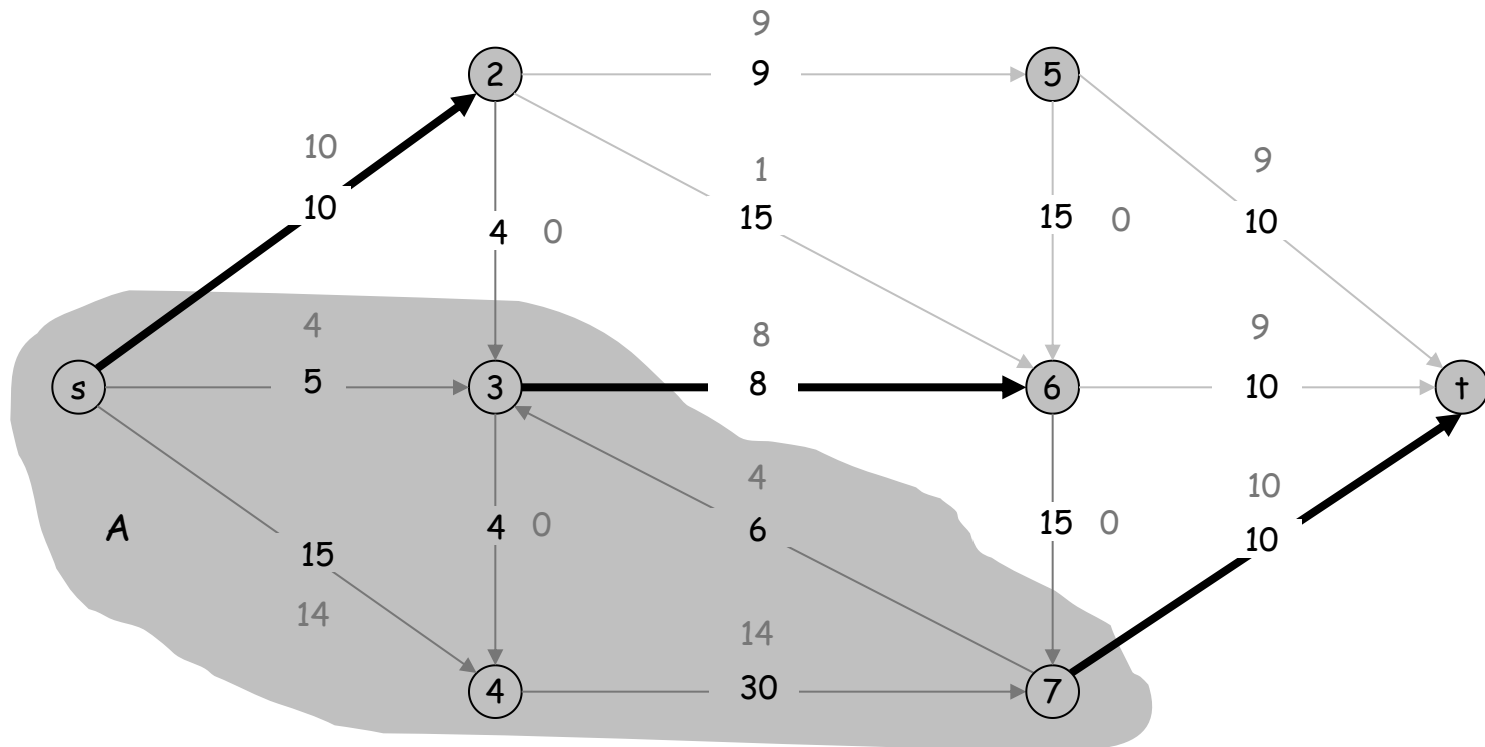


Certificate of Optimality

Corollary. Let f be any flow, and let (A, B) be any cut. If $v(f) = \text{cap}(A, B)$, then f is a max flow and (A, B) is a min cut.

Value of flow = 28


Cut capacity = 28 \Rightarrow Flow value ≤ 28



Towards a Max Flow Algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for all edge $e \in E$.
- Find an s-t path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.




Flow value = 0

Towards a Max Flow Algorithm

Greedy algorithm.

- Start with $f(e) = 0$ for all edge $e \in E$.
- Find an s-t path P where each edge has $f(e) < c(e)$.
- Augment flow along path P .
- Repeat until you get stuck.

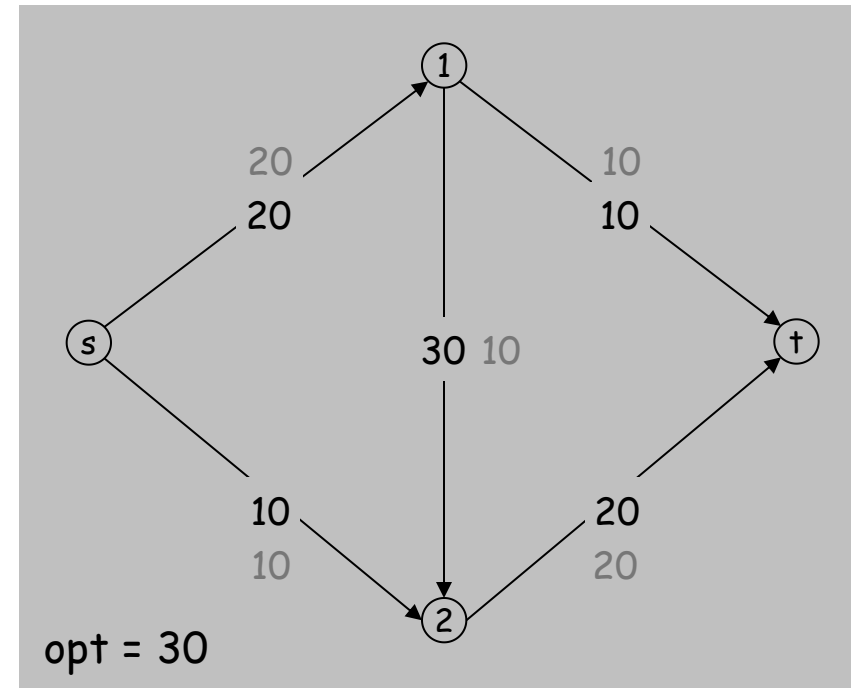
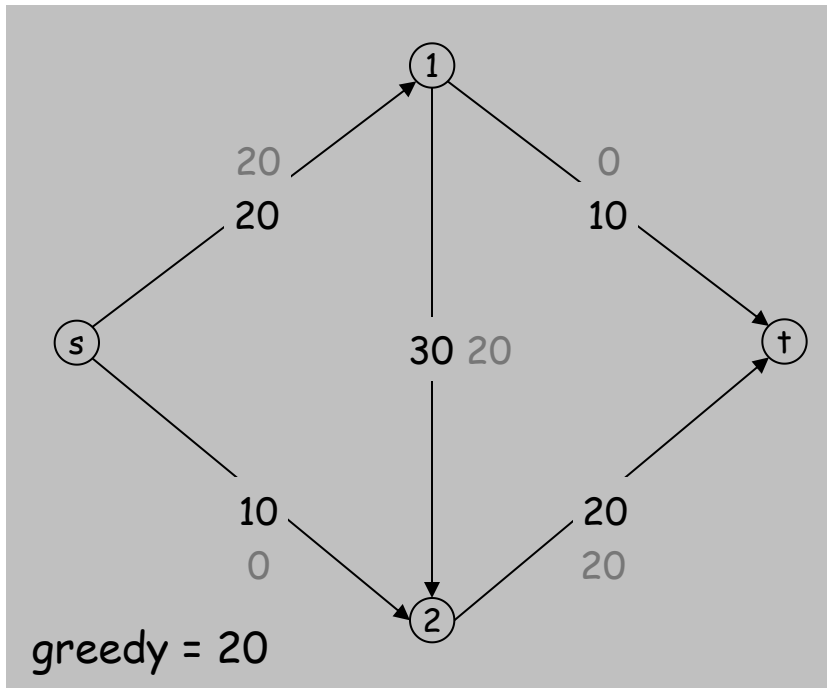


Flow value = 20

Towards a Max Flow Algorithm

Greedy algorithm.

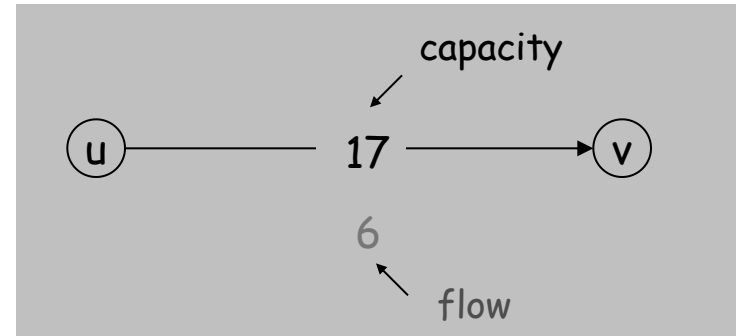
- Start with $f(e) = 0$ for all edge $e \in E$.
 - Find an s-t path P where each edge has $f(e) < c(e)$.
 - Augment flow along path P .
 - Repeat until you get **stuck**.
- ↖ locally optimality \nRightarrow global optimality



Residual Graph

Original edge: $e = (u, v) \in E$.

- Flow $f(e)$, capacity $c(e)$.

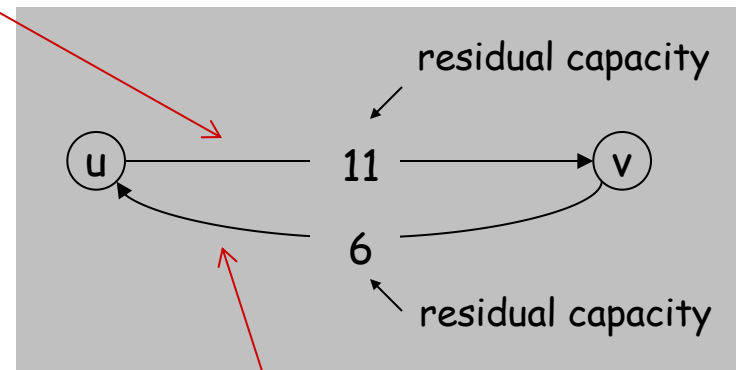


Residual edge.

- "Undo" flow sent.
- $e = (u, v)$ and $e^R = (v, u)$.
- Residual capacity:

$$c_f(e) = \begin{cases} c(e) - f(e) & \text{if } e \in E \\ f(e) & \text{if } e^R \in E \end{cases}$$

Forward edge



Backward edge

Residual graph: $G_f = (V, E_f)$.

- Residual edges with positive residual capacity.
- $E_f = \{e : f(e) < c(e)\} \cup \{e^R : f(e) > 0\}$.

Augmenting path

Def. An **augmenting path** is a simple $s \rightarrow t$ path in the residual graph G_f

Def. The **bottleneck capacity** of an augmenting path P is the minimum residual capacity of any edge in P .

Key property. Let f be a flow and let P be an augmenting path in G_f , then after calling $f' \leftarrow \text{Augment}(f, c, P)$, the resulting f' is flow and

$$v(f') = v(f) + \text{bottleneck}(G_f, P)$$

Augmenting Path Algorithm

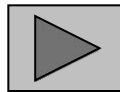
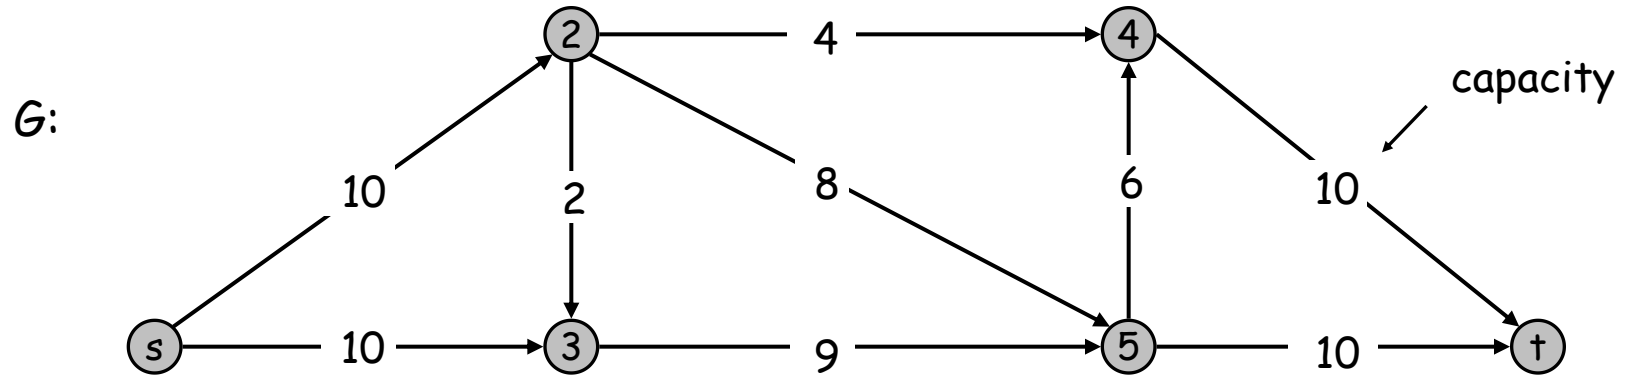
```
Augment(f, c, P) {  
    b ← bottleneck(P)  
    foreach e ∈ P {  
        if (e ∈ E) f(e) ← f(e) + b  
        else      f(eR) ← f(eR) - b  
    }  
    return f  
}
```

forward edge

reverse edge

```
Ford-Fulkerson(G, s, t, c) {  
    foreach e ∈ E f(e) ← 0  
    Gf ← residual graph  
  
    while (there exists augmenting path P) {  
        f ← Augment(f, c, P)  
        update Gf  
    }  
    return f  
}
```

Ford-Fulkerson Algorithm



Max-Flow Min-Cut Theorem

Augmenting path theorem. Flow f is a max flow iff there are no augmenting paths.

Max-flow min-cut theorem. [Elias-Feinstein-Shannon 1956, Ford-Fulkerson 1956]
The value of the max flow is equal to the value of the min cut.

Pf. We prove both simultaneously by showing TFAE (the following are equivalent) :

- (i) There exists a cut (A, B) such that $v(f) = \text{cap}(A, B)$.
- (ii) Flow f is a max flow.
- (iii) There is no augmenting path relative to f .

(i) \Rightarrow (ii) This was the corollary to weak duality lemma. (Slide 17)

(ii) \Rightarrow (iii) We show contrapositive.

- Let f be a max flow. If there exists an augmenting path, then we can improve f by sending flow along path.

Proof of Max-Flow Min-Cut Theorem


(iii) \Rightarrow (i)

- Let f be a flow with no augmenting paths.
- Let A be set of vertices reachable from s in residual graph.
- By definition of A , $s \in A$.
- By definition of f , $t \notin A$.

$f(e) = 0$, if not, there will be a backward edge in G_f ,
Violate no augmenting paths in G_f

$$\begin{aligned} v(f) &= \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) \\ &= \sum_{e \text{ out of } A} c(e) \\ &= \text{cap}(A, B) \quad \blacksquare \end{aligned}$$

If not, there will be a forward edge in G_f ,
Violate no augmenting paths in G_f



original network

Running Time

Assumption. All capacities are integers between 1 and $C = \sum_{e \text{ out of } s} c(e)$.

Invariant. Every flow value $f(e)$ and every residual capacity $c_f(e)$ remains an integer throughout the algorithm.

Theorem. The algorithm terminates in at most $v(f^*) \leq C$ iterations.

Pf. Each augmentation increase value by at least 1. ■

To find an s - t path in G_f , say by BFS, $O(m+n)$ with $m \geq n/2$,
Procedure $\text{augment}(f, P)$ takes $O(n)$, as the path has at most $n-1$ edges


Corollary. If $C = 1$, Ford-Fulkerson runs in $O(mn)$ time. ↗

Integrality theorem. If all capacities are integers, then there exists a max flow f for which every flow value $f(e)$ is an integer.

Pf. Since algorithm terminates, theorem follows from invariant. ■

Chapter 7

Network Flow



Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.

7.3 Choosing Good Augmenting Paths

Choosing good augmenting paths


Use care when selecting augmenting paths

- Some choices lead to exponential algorithms
- Clever choice lead to polynomial algorithms


Pathology. When edge capacities can be irrational, no guarantee that Ford-Fulkerson terminates (or converges to a maximum flow)!

Goal. Choose augmenting paths so that:

- Can find augmenting paths efficiently.
- Few iterations




Forward edge



Backward edge

G_f :






Figure Parts (a) through (d) depict four iterations of the Ford-Fulkerson Algorithm using a bad choice of augmenting paths: The augmentations alternate between the path P_1 through the nodes s, u, v, t in order and the path P_2 through the nodes s, v, u, t in order.

Choosing good augmenting paths


Choose augmenting paths with:

- Max bottleneck capacity ("fattest"). ← how to find?
- Sufficiently large bottleneck capacity. ← next
- Fewest edges. ← ahead

Capacity-scaling algorithm

Overview. Choosing augmenting paths with “large” bottleneck capacity.

- Maintain scaling parameter Δ .
- Let $G_f(\Delta)$ be the part of the residual graph containing only those edges with capacity $\geq \Delta$.
- Any augmenting path in $G_f(\Delta)$ has bottleneck capacity $\geq \Delta$.



Scaling Max-Flow

Initially $f(e)=0$ for all e in G

Initially set Δ to be the largest power of 2 that is no larger than the maximum capacity out of s : $\Delta \leq \max_{e \text{ out of } s} c_e$

While $\Delta \geq 1$

While there is an s - t path in the graph $G_f(\Delta)$

Let P be a simple s - t path in $G_f(\Delta)$

$f' = \text{augment}(f, P)$

Update f to be f' and update $G_f(\Delta)$

Endwhile

$\Delta = \Delta/2$

Endwhile

Return f

Capacity-scaling algorithm: proof of correctness

Assumption: All edge capacities are integers between 1 and C .

Invariant. The scaling parameter Δ is a power of 2.

Pf. Initially a power of 2 (largest power of $2 \leq C$); each phase divides Δ by exactly 2.

Integrality invariant. Throughout the algorithm, every edge flow $f(e)$ and residual capacity $c_f(e)$ is an integer.

Pf. Same as for genetic Ford-Fulkerson.

Theorem. If capacity-scaling algorithm terminates, then f is a max flow.

Pf.

- By integrality invariant, when $\Delta = 1 \Rightarrow G_f(\Delta) = G_f$.
- Upon termination of $\Delta = 1$ phase, there are no augmenting paths.
- Result follows augmenting path theorem.

Capacity-scaling algorithm: analysis of running time

Lemma 1. There are $1 + \lceil \log_2 C \rceil$ scaling phases.

Pf. Initial $C/2 < \Delta \leq C$; Δ decreases by a factor of 2 in each iteration.

Lemma 2. Let f be the flow at the end of a Δ -scaling phase, then the max-flow value $\leq v(f) + m\Delta$.


Pf. Next slide.

Lemma 3. There are $\leq 2m$ augmentations per scaling phase.

Pf.

- Let f be the flow at the beginning of a Δ -scaling phase.
- Lemma 2 \Rightarrow max-flow value $\leq v(f) + m(2\Delta)$.
- Each augmentation in a Δ -scaling phase increases $v(f)$ by at least Δ .

or equivalently,
at the end
of a 2Δ -scaling phase



Theorem. The capacity-scaling algorithm takes $O(m^2 \log C)$ time.

Pf.

- Lemma 1 + Lemma 3 $\Rightarrow O(m \log C)$ augmentations.
- Finding an augmenting path takes $O(m)$ time.


Capacity-scaling algorithm: analysis of running time

Lemma 2. Let f be the flow at the end of a Δ -scaling phase, then the max-flow value $\leq v(f) + m\Delta$.

Pf.

- We show there exists a cut (A, B) such that $\text{cap}(A, B) \leq v(f) + m\Delta$.
- Choose A to be the set of nodes reachable from s in $G_f(\Delta)$.
- By definition of flow f : $t \notin A$.

$$\begin{aligned}
 \text{flow value lemma} \quad val(f) &= \sum_{e \text{ out of } A} f(e) - \sum_{e \text{ in to } A} f(e) \\
 &\geq \sum_{e \text{ out of } A} (c(e) - \Delta) - \sum_{e \text{ in to } A} \Delta \\
 &\geq \sum_{e \text{ out of } A} c(e) - \sum_{e \text{ out of } A} \Delta - \sum_{e \text{ in to } A} \Delta \\
 &\geq \text{cap}(A, B) - m\Delta \quad \blacksquare
 \end{aligned}$$



Residual capacity:

$$c_f(e) = \begin{cases} c(e) - f(e) & \text{if } e \in E \\ f(e) & \text{if } e^R \in E \end{cases}$$

Shortest augmenting path

Q. How to choose next augmenting path in Ford-Fulkerson?

A. Pick one that uses the fewest edges.

can find via BFS

SHORTEST-AUGMENTING-PATH(G)

FOREACH $e \in E : f(e) \leftarrow 0$.

$G_f \leftarrow$ residual network of G with respect to flow f .

WHILE (there exists an $s \rightarrow t$ path in G_f)

$P \leftarrow \text{BREADTH-FIRST-SEARCH}(G_f)$.


$f \leftarrow \text{AUGMENT}(f, c, P)$.

Update G_f .

RETURN f .

Chapter 7

Network Flow




Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.

7.5 Bipartite Matching

Matching

Matching.

- Input: undirected graph $G = (V, E)$.
- $M \subseteq E$ is a **matching** if each node appears in at most one edge in M .
- Max matching: find a max cardinality matching.




Bipartite Matching

Bipartite matching.

- Input: undirected, **bipartite** graph $G = (L \cup R, E)$.
- $M \subseteq E$ is a **matching** if each node appears in at most one edge in M .
- Max matching: find a max cardinality matching.

Def. A graph G is **bipartite** if the nodes can be partitioned into two subsets L and R such that every edge connects a node in L with a node in R .




matching
1-2', 3-1', 4-5'

Bipartite Matching

Bipartite matching.

- Input: undirected, **bipartite** graph $G = (L \cup R, E)$.
- $M \subseteq E$ is a **matching** if each node appears in at most one edge in M .
- Max matching: find a max cardinality matching.




max matching
1-1', 2-2', 3-3' 4-4'

Bipartite Matching

Max flow formulation.

- Create digraph $G' = (L \cup R \cup \{s, t\}, E')$.
 - Direct all edges from L to R , and assign infinite (or unit) capacity.
 - Add source s , and unit capacity edges from s to each node in L .
 - Add sink t , and unit capacity edges from each node in R to t .





Bipartite Matching: Proof of Correctness

Theorem. value of max flow in G' = Max cardinality matching in G .

Pf. \leq

- Given max matching M of cardinality k .
- Consider flow f that sends 1 unit along each of k paths.
- f is a flow, and has cardinality k . ▪





Bipartite Matching: Proof of Correctness

Theorem. value of max flow in G' = Max cardinality matching in G .

Pf. \geq

- Let f be a max flow in G' of value k .
- Integrality theorem $\Rightarrow k$ is integral and can assume f is 0-1.
- Consider M = set of edges from L to R with $f(e) = 1$.
 - each node in L and R participates in at most one edge in M
 - $|M| = k$: apply **flow-value lemma** to cut $(L \cup s, R \cup t)$ ■

Flow value lemma. Let f be any flow, and let (A, B) be any s - t cut. Then, the net flow sent across the **cut** is equal to the amount leaving s .



Perfect Matching

Def. A matching $M \subseteq E$ is **perfect** if each node appears in exactly one edge in M .

Q. When does a bipartite graph have a perfect matching?

Structure of bipartite graphs with perfect matchings.

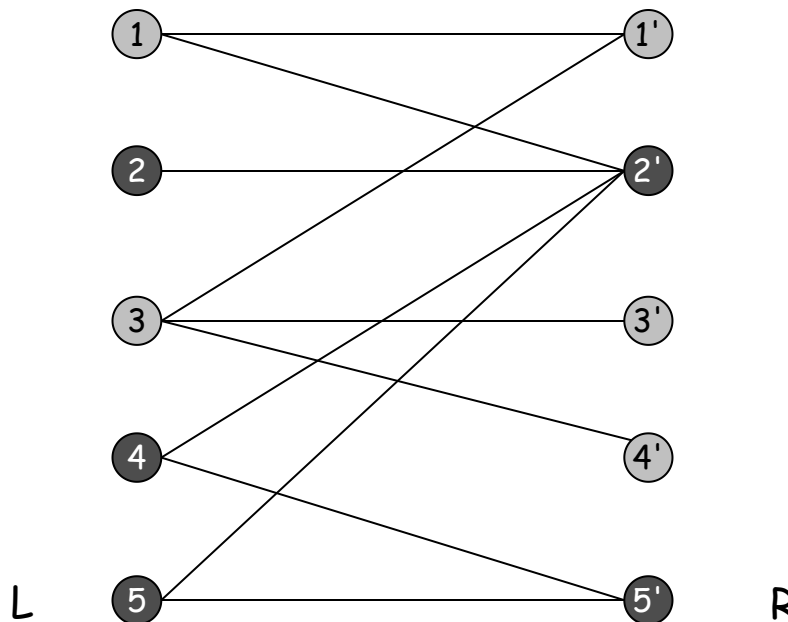
- Clearly we must have $|L| = |R|$.
- What other conditions are necessary?
- What conditions are sufficient?

Perfect Matching

Notation. Let S be a subset of nodes, and let $N(S)$ be the set of nodes adjacent to nodes in S .

Observation. If a bipartite graph $G = (L \cup R, E)$, has a perfect matching, then $|N(S)| \geq |S|$ for all subsets $S \subseteq L$.

Pf. Each node in S has to be matched to a different node in $N(S)$.



No perfect matching:


$S = \{ 2, 4, 5 \}$

$N(S) = \{ 2', 5' \}$.

Hall's marriage theorem

Theorem. [Frobenius 1917, Hall 1935] Let $G = (L \cup R, E)$ be a bipartite graph with $|L| = |R|$. Then, graph G has a perfect matching iff $|N(S)| \geq |S|$ for all subsets $S \subseteq L$.

Pf. \Rightarrow This was the previous observation.



No perfect matching:

$S = \{ 2, 4, 5 \}$

$N(S) = \{ 2', 5' \}$.

Bipartite Matching: Running Time

Which max flow algorithm to use for bipartite matching?


- Generic augmenting path: $O(mn \text{ val}(f^*)) = O(mnC)$.
- Capacity scaling: $O(m^2 \log C)$.

Non-bipartite matching.

- Structure of non-bipartite graphs is more complicated, but well-understood. [Tutte-Berge, Edmonds-Galai]
- Blossom algorithm: $O(n^4)$. [Edmonds 1965]
- Best known: $O(m n^{1/2})$. [Micali-Vazirani 1980]

Chapter 7

Network Flow



Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.


7.6 Disjoint Paths in Directed and Undirected Graphs

Edge-disjoint paths

Def. Two paths are **edge-disjoint** if they have no edge in common.

Edge-disjoint paths problem. Given a digraph $G = (V, E)$ and two nodes s and t , find the max number of edge-disjoint $s \rightsquigarrow t$ paths.

Ex. Communication networks.




Edge-disjoint paths

Def. Two paths are **edge-disjoint** if they have no edge in common.

Edge-disjoint paths problem. Given a digraph $G = (V, E)$ and two nodes s and t , find the max number of edge-disjoint $s \rightsquigarrow t$ paths.

Ex. Communication networks.




Edge-disjoint paths

Max-flow formulation. Assign unit capacity to every edge.

Theorem. Max number of edge-disjoint $s \rightsquigarrow t$ paths = value of max flow.

Pf. \geq

- Suppose there are k edge-disjoint $s \rightsquigarrow t$ paths P_1, \dots, P_k .
- Set $f(e) = 1$ if e participates in some path P_j ; else set $f(e) = 0$.
- Since paths are edge-disjoint, f is a flow of value k . ▀




Edge-disjoint paths

Max-flow formulation. Assign unit capacity to every edge.

Theorem. Max number of edge-disjoint $s \rightsquigarrow t$ paths = value of max flow.

Pf. \leq

- Suppose max flow value is k .
- Integrality theorem \Rightarrow there exists 0-1 flow f of value k .
- Consider edge (s, u) with $f(s, u) = 1$.
 - by flow conservation, there exists an edge (u, v) with $f(u, v) = 1$
 - continue until reach t , always choosing a new edge
- Produces k (not necessarily simple) edge-disjoint paths




can eliminate cycles
to get simple paths
in $O(mn)$ time if desired
(flow decomposition)

Network connectivity

Def. A set of edges $F \subseteq E$ **disconnects** t from s if every $s \rightsquigarrow t$ path uses at least one edge in F .

Network connectivity. Given a digraph $G = (V, E)$ and two nodes s and t , find min number of edges whose removal disconnects t from s .

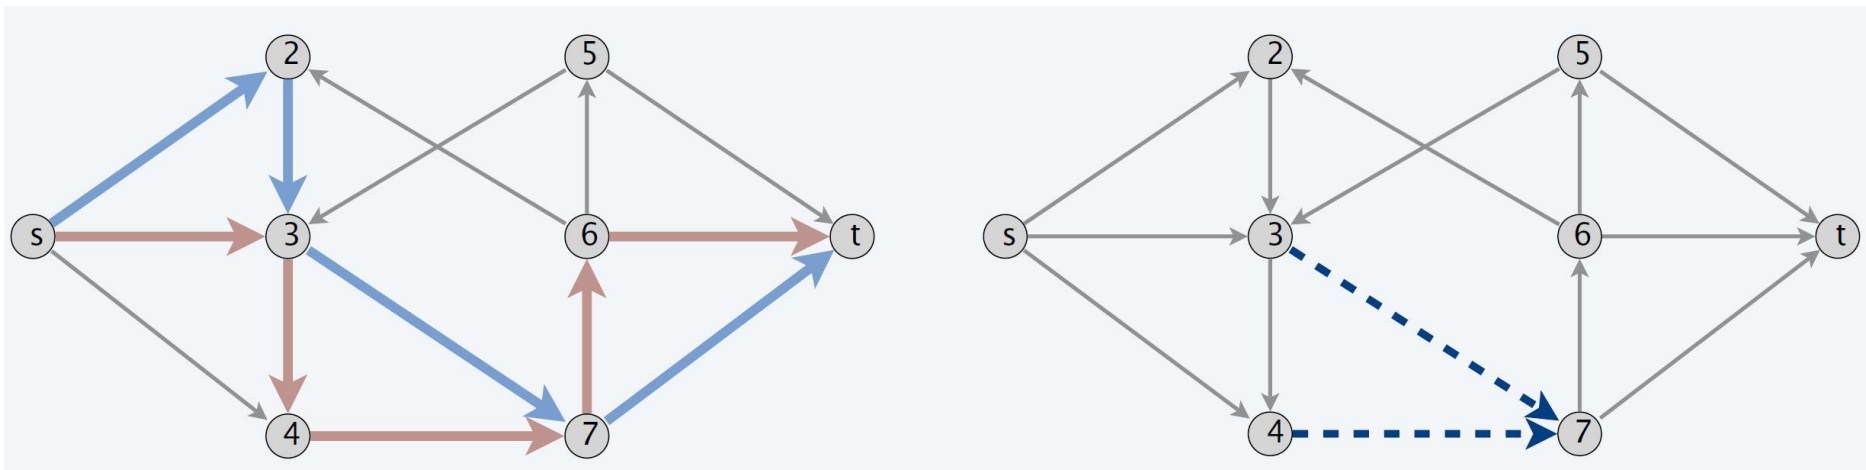


Menger's theorem

Theorem. [Menger 1927] The max number of edge-disjoint $s \rightsquigarrow t$ paths equals the min number of edges whose removal disconnects t from s .

Pf. \leq

- Suppose the removal of $F \subseteq E$ disconnects t from s , and $|F| = k$.
- Every $s \rightsquigarrow t$ path uses at least one edge in F .
- Hence, the number of edge-disjoint paths is $\leq k$.




Menger's theorem

Theorem. [Menger 1927] The max number of edge-disjoint $s \rightsquigarrow t$ paths equals the min number of edges whose removal disconnects t from s .

Pf. \geq


- Suppose max number of edge-disjoint paths is k .
- Then value of max flow = k .
- Max-flow min-cut theorem \Rightarrow there exists a cut (A, B) of capacity k .
- Let F be set of edges going from A to B .
- $|F| = k$ and disconnects t from s .



Edge-disjoint paths in undirected graphs

Def. Two paths are **edge-disjoint** if they have no edge in common.


Edge-disjoint paths problem in undirected graphs. Given a graph $G = (V, E)$ and two nodes s and t , find the max number of edge-disjoint s - t paths.



Edge-disjoint paths in undirected graphs

Def. Two paths are **edge-disjoint** if they have no edge in common.


Edge-disjoint paths problem in undirected graphs. Given a graph $G = (V, E)$ and two nodes s and t , find the max number of edge-disjoint s - t paths.



Edge-disjoint paths in undirected graphs

Def. Two paths are **edge-disjoint** if they have no edge in common.

Edge-disjoint paths problem in undirected graphs. Given a graph $G = (V, E)$ and two nodes s and t , find the max number of edge-disjoint s - t paths.




Edge-disjoint paths in undirected graphs

Max-flow formulation. Replace each edge with two antiparallel edges and assign unit capacity to every edge.

Observation. Two paths P_1 and P_2 may be edge-disjoint in the digraph but not edge-disjoint in the undirected graph.

if P_1 uses edge (u, v)
and P_2 uses its antiparallel edge (v, u)




Edge-disjoint paths in undirected graphs

Max-flow formulation. Replace each edge with two antiparallel edges and assign unit capacity to every edge.

Lemma. In any flow network, there exists a maximum flow f in which for each pair of antiparallel edges e and e' : either $f(e) = 0$ or $f(e') = 0$ or both. Moreover, integrality theorem still holds.

Pf. [by induction on number of such pairs]

- Suppose $f(e) > 0$ and $f(e') > 0$ for a pair of antiparallel edges e and e' .
- Set $f(e) = f(e) - \delta$ and $f(e') = f(e') - \delta$, where $\delta = \min \{ f(e), f(e') \}$.
- f is still a flow of the same value but has one fewer such pair




Edge-disjoint paths in undirected graphs

Max-flow formulation. Replace each edge with two antiparallel edges and assign unit capacity to every edge.

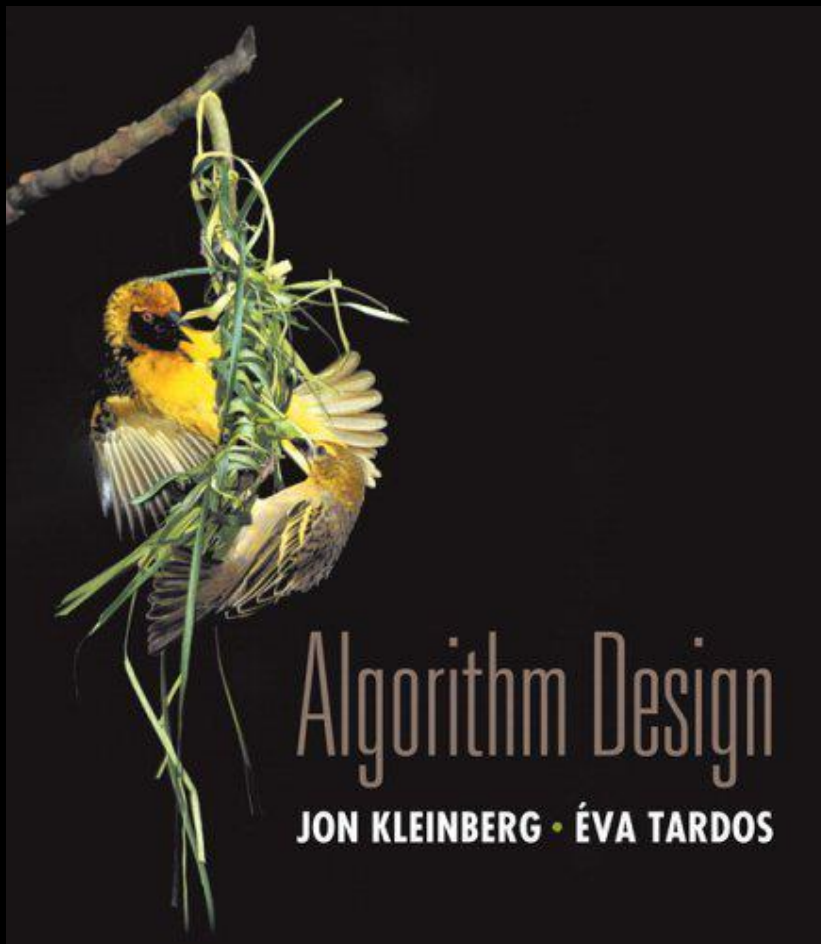
Lemma. In any flow network, there exists a maximum flow f in which for each pair of antiparallel edges e and e' : either $f(e) = 0$ or $f(e') = 0$ or both. Moreover, integrality theorem still holds.

Theorem. Max number of edge-disjoint $s \rightsquigarrow t$ paths = value of max flow.
Pf. Similar to proof in digraphs; use lemma.



Chapter 7

Network Flow




Slides by Kevin Wayne.
Copyright © 2005 Pearson-Addison Wesley.
All rights reserved.

7.7 Extensions to Maximum-Flow Problem

Multiple sources and sinks

Def. Given a digraph $G = (V, E)$ with edge capacities $c(e) \geq 0$ and multiple source nodes and multiple sink nodes, find max flow that can be sent from the source nodes to the sink nodes.

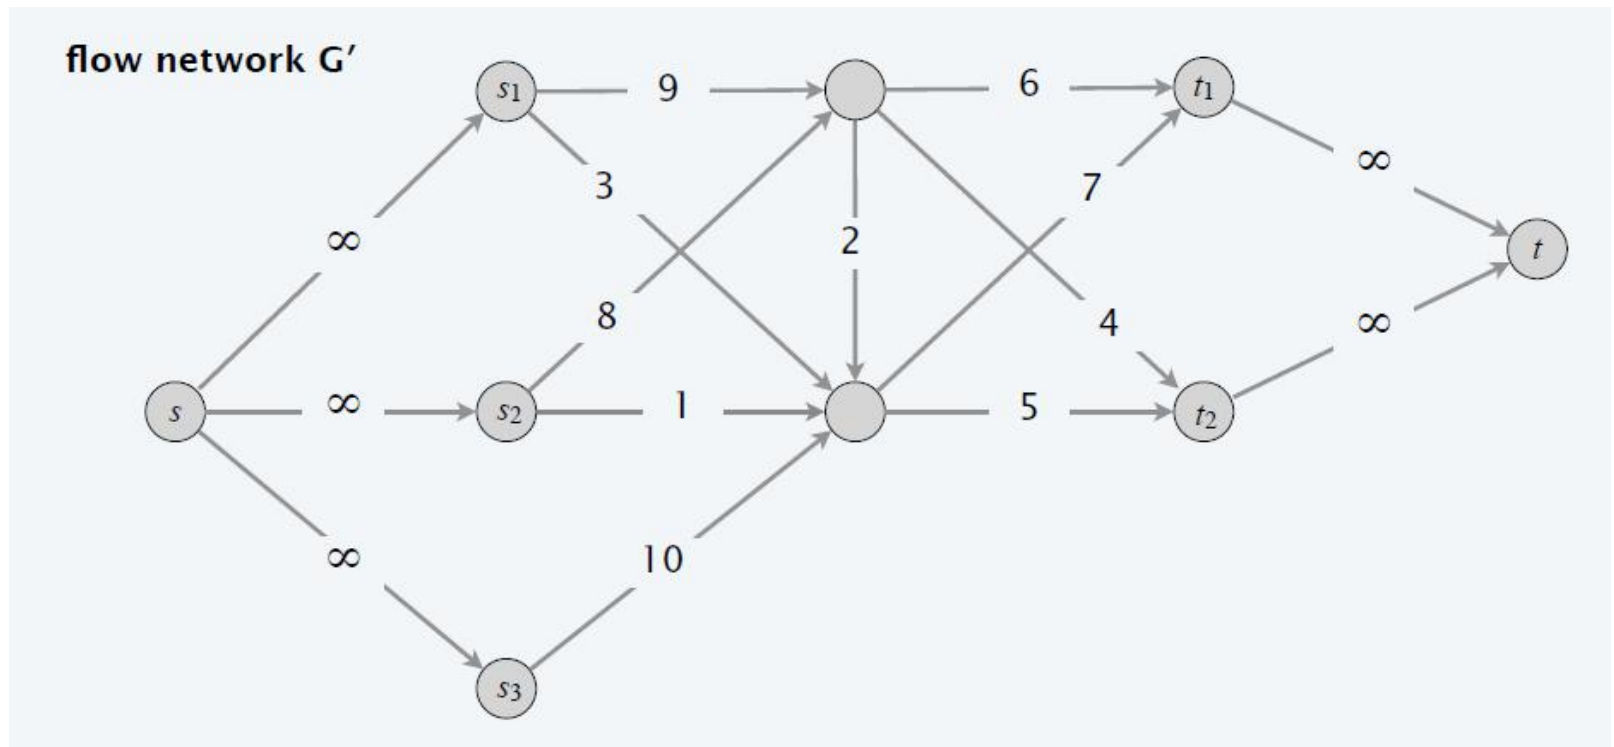
flow network G



Multiple sources and sinks: max-flow formulation

- Add a new source node s and sink node t .
- For each original source node s_i add edge (s, s_i) with capacity ∞ .
- For each original sink node t_j , add edge (t_j, t) with capacity ∞ .

Claim. 1-1 correspondence between flows in G and G' .




Circulation with supplies and demands

Def. Given a digraph $G = (V, E)$ with edge capacities $c(e) \geq 0$ and node demands $d(v)$, a **circulation** is a function $f(e)$ that satisfies:

For each $e \in E$: $0 \leq f(e) \leq c(e)$ (capacity)

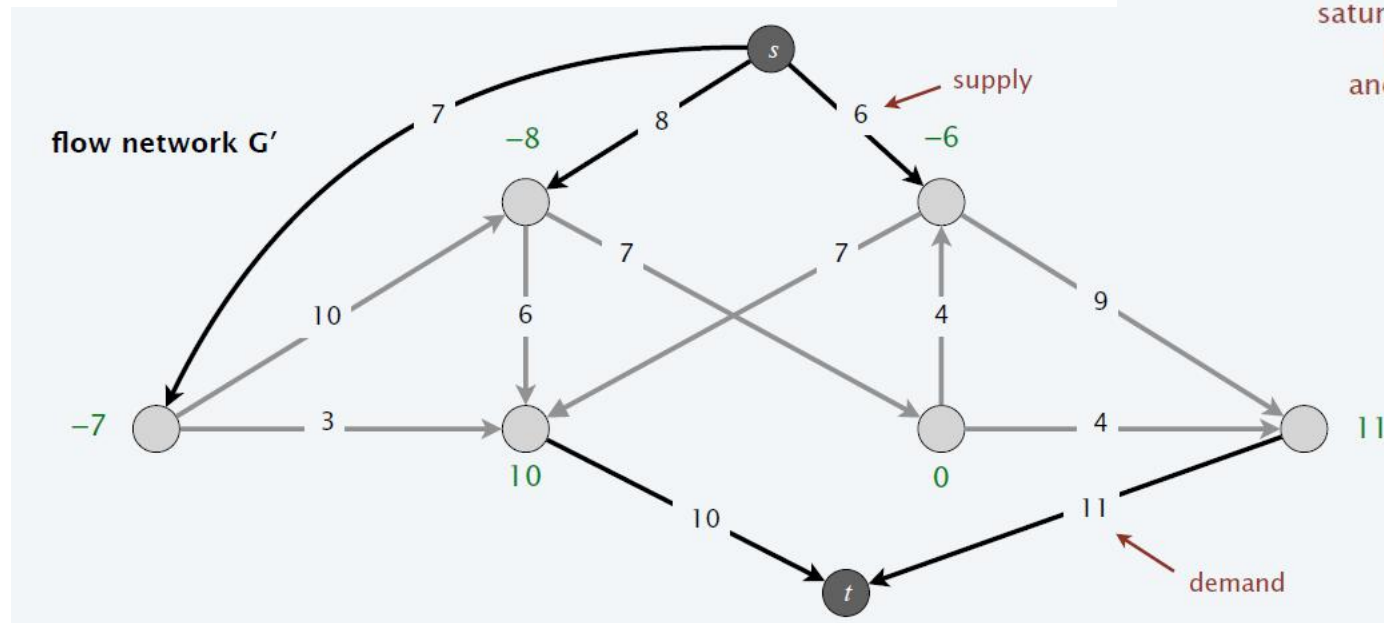
For each $v \in V$: $\sum_{e \text{ in to } v} f(e) - \sum_{e \text{ out of } v} f(e) = d(v)$ (flow conservation)



Circulation with supplies and demands: max-flow formulation

- Add new source s and sink t .
- For each v with $d(v) < 0$, add edge (s, v) with capacity $-d(v)$.
- For each v with $d(v) > 0$, add edge (v, t) with capacity $d(v)$.

Claim. G has circulation iff G' has max flow of value $D = \sum_{v: d(v) > 0} d(v) = \sum_{v: d(v) < 0} -d(v)$



↑
saturates all edges
leaving s
and entering t

Circulation with supplies and demands

Integrality theorem. If all capacities and demands are integers, and there exists a circulation, then there exists one that is integer-valued.

Pf. Follows from max-flow formulation + integrality theorem for max flow.

Theorem. Given (V, E, c, d) , there does **not** exist a circulation iff there exists a node partition (A, B) such that $\sum_{v \in B} d(v) > \text{cap}(A, B)$.

Pf sketch. Look at min cut in G' .

demand by nodes in B exceeds
supply of nodes in B plus
max capacity of edges going from A to B

Previous slide: G has circulation iff G' has max flow of value \Rightarrow max flow \Rightarrow min cut